

Raj Patel

East Hartford, CT | (860)-709-1318 | rajpat243@gmail.com | [linkedin.com/in/rajpat243](https://www.linkedin.com/in/rajpat243) | github.com/rajpat243

SKILLS

Programming:	Python (Pandas, NumPy, Matplotlib), SQL (MySQL, Postgres), NoSQL (MongoDB), JavaScript
Frameworks:	dbt, Spark, Qlik Sense, React Js, Flask
Cloud & Tools:	AWS (Glue, Lambda, S3, Bedrock), Snowflake, Databricks, Azure AD, GitHub/GitLab, VS Code
Methodologies:	ETL/ELT Pipeline Design, Data Architecture, Agile

EXPERIENCE

Data Engineer , Travelers Insurance	Jun 2025 - Present Hartford, CT
- Architected and owned production ETL/ELT pipelines in dbt that ingest and transform data into Qlik Sense applications, enabling business partners to onboard new features and data sources with minimal engineering support. - Built and deployed a custom Python MCP server connecting Claude AI to on-premises Qlik Sense Enterprise data, enabling natural language querying without Qlik Cloud dependency and eliminating analyst reliance on ad-hoc engineering requests. - Partnered with business stakeholders and cross-functional teams to define data requirements, translate them into data models, and deliver scalable data architecture aligned with governance and security policies. - Built a Claude AI skill to accelerate source-to-target mapping workflows, auto-generating Snowflake field discovery queries and structured CSV mapping files, reducing manual STTM effort for the data engineering team.	
Embedded Software Engineering Intern , General Dynamics Electric Boat	Jun 2024 - Jan 2025 Groton, CT
- Developed and tested C++ code to validate the functionality of the control panel on Columbia-class submarines, applying structured unit test plans to ensure reliability and compliance with system requirements. - Wrote C++ code to configure specific controls within the Columbia-class control panel, ensuring accurate control behavior and compliance with embedded system requirements. - Built a Python-based automated test harness that ingested and validated embedded firmware output against requirement datasets, cutting regression test time from 4 hours to 30 minutes.	

PROJECTS

AI-Powered Onboarding Assistant (Hackathon)	Sept 2025
- Designed and built an AI onboarding platform leveraging prior employee data to recommend role-appropriate permissions and SailPoint identity groups, reducing manual provisioning overhead. - Trained an AWS Bedrock Agent Core agent on synthetic employee and permissions data, enabling natural language Q&A for new hires and employees switching teams — surfacing accurate access policies without IT escalation. - Built structured data pipelines to ingest and query employee identity records as the AI agent's knowledge base.	
Real-Time Shark Tracking and Classification System	Aug 2024 - May 2025
- Engineered data pipelines using the ClusterDuck Protocol to transmit real-time sensor and movement data from maritime environments into OWL's cloud-based Data Management System. - Built and trained a CNN machine learning model to classify sharks vs. non-shark aquatic animals, achieving 85% accuracy on shark images and 80% overall accuracy. - Optimized data processing for edge deployment on the Coral Dev Board, enabling real-time ML inferencing and efficient handling of large streaming datasets.	

CERTIFICATIONS

- AWS Cloud Practitioner | Issued Nov 2025
- AWS AI Practitioner | Issued Jan 2026
- AWS Data Engineer – Associate | Issued May 2026
- AWS Developer – Associate | Issued May 2026

EDUCATION

University of Connecticut - Storrs	May 2025
<i>B.S. in Computer Science, Concentration in Software Engineering and Development</i>	<i>GPA: 3.70/4.00</i>
Relevant Coursework: Artificial Intelligence, Data Structures and Algorithms, Object Oriented Programming, Principles of Databases, Cloud Computing	